

Project Three for Data Mining

Raymond Anthony Ford
raford2@miners.utep.edu

Due: 06 March 2018

Contents

1	Theoretical problem	1
2	Computer project	2
2.1	Bringing the data in	2
2.2	Data cleaning	4
2.3	Obtaining a distance matrix	5
2.4	Clustering algorithms	6
2.4.1	K -means clustering	6
2.4.2	Partition around medoids (PAM) clustering	8
2.4.3	Comparing K -means and PAM	10
2.5	<i>Post hoc</i> analysis	10
3	Discussion	12
4	References	12

1 Theoretical problem

Let \mathcal{C} denote a clustering scheme that clusters data into K clusters $\{c_1, \dots, c_K\}$. The within-cluster point scatter $W(\mathcal{C})$ is defined as

$$W(\mathcal{C}) = \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \|\mathbf{x}_i - \mathbf{x}_{i'}\|^2.$$

We need to show that

$$W(\mathcal{C}) = \sum_{k=1}^K n_k \sum_{i \in c_k} \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2,$$

where $\bar{\mathbf{x}}_k$ denotes the mean vector and n_k is the sample size for cluster c_k .

Proof: We begin this proof by observing the following.

$$\begin{aligned}
W(\mathcal{C}) &= \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \|\mathbf{x}_i - \mathbf{x}_{i'}\|^2 \\
&= \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \|\mathbf{x}_i - \bar{\mathbf{x}}_k + \bar{\mathbf{x}}_k - \mathbf{x}_{i'}\|^2 \\
&= \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \{\|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2 + \|\bar{\mathbf{x}}_k - \mathbf{x}_{i'}\|^2 + 2(\mathbf{x}_i - \bar{\mathbf{x}}_k)^T(\mathbf{x}_{i'} - \bar{\mathbf{x}}_k)\}
\end{aligned}$$

Moreover, note that:

$$2 \sum_{i' \in c_k} (\mathbf{x}_i - \bar{\mathbf{x}}_k)^T (\mathbf{x}_{i'} - \bar{\mathbf{x}}_k) = 2(\mathbf{x}_i - \bar{\mathbf{x}}_k)^T \sum_{i' \in c_k} (\mathbf{x}_{i'} - \bar{\mathbf{x}}_k) \propto \sum_{i' \in c_k} (\mathbf{x}_{i'} - \bar{\mathbf{x}}_k) = \sum_{i' \in c_k} \mathbf{x}_{i'} - \sum_{i' \in c_k} \bar{\mathbf{x}}_k = 0.$$

Thus we have,

$$\begin{aligned}
W(\mathcal{C}) &= \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2 + \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \|\bar{\mathbf{x}}_k - \mathbf{x}_{i'}\|^2 + 0 \\
&= \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} n_k \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2 + \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2 \\
&= \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} n_k \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2 + \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} n_k \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2 \\
&= \sum_{k=1}^K n_k \sum_{i \in c_k} \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2.
\end{aligned}$$

Consequently, we have $W(\mathcal{C}) = \frac{1}{2} \sum_{k=1}^K \sum_{i \in c_k} \sum_{i' \in c_k} \|\mathbf{x}_i - \mathbf{x}_{i'}\|^2 = \sum_{k=1}^K n_k \sum_{i \in c_k} \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|^2$, as was to be shown.

2 Computer project

2.1 Bringing the data in

We begin our venture for this project by first bringing the data into R, and then by listing the missing rate (in terms of percentage) for each of the variables in our data set with the following code.

```
dat <- read.csv('~/.Dropbox/Spring2018/STAT5474/HMEQ.csv', header=TRUE, sep=',')
vnames <- colnames(dat)
```

```

n <- nrow(dat)
out <- NULL
for (j in 1:ncol(dat)){
  vname <- colnames(dat)[j]
  x <- as.vector(dat[,j])
  nmiss <- sum(is.na(x))
  ncomplete <- n-nmiss
  out <- rbind(out, c(column.number=j, variable.name=vname,
                      missing.percentage=(nmiss/n)*100))
}
out <- as.data.frame(out)
row.names(out) <- NULL
out

```

```

##      column.number variable.name missing.percentage
## 1              1          BAD              0
## 2              2          LOAN              0
## 3              3      MORTDUE      8.69127516778524
## 4              4          VALUE      1.87919463087248
## 5              5      REASON      4.22818791946309
## 6              6          JOB      4.68120805369128
## 7              7          YOJ      8.64093959731544
## 8              8      DEROG      11.8791946308725
## 9              9      DELINQ      9.73154362416107
## 10             10      CLAGE      5.16778523489933
## 11             11      NINQ      8.55704697986577
## 12             12      CLNO      3.7248322147651
## 13             13      DEBTINC      21.258389261745

```

We see that we have missing values for all of the variables except BAD and LOAN, so we will need to do some missing value imputation. The good news is that we do not have any missing values for our target variable BAD.

We can visualize the missing values in our data set in an attempt to try to identify any potential patterns in the missing values.

```

suppressMessages(library(VIM, quietly=TRUE))
aggr(dat, col=c('blue3', 'darkgoldenrod'), numbers=TRUE, labels=names(dat),
      cex.axis=.7, gap=5, ylab=c("Missing data", "Pattern"))

```

```

## Warning in plot.aggr(res, ...): not enough vertical space to display
## frequencies (too many combinations)

```

We get the message about not having enough vertical space because of the size of our data, but the subset of the data that is displayed in Figure 1 is still informative. We see that for many of the applications there are multiple missing values for each of the variables. In fact, for at least two applicants, almost all of the predictive variables are missing except for a




Figure 1: Plots detailing the missing values in our dataset. The plot on the left shows the proportion of missing values for each variable, and the plot on the right shows the pattern for the missing values with all of the variables.

few. This seems to be pretty common for the subset of the data that is displayed. This is an intriguing finding given that banks have become more conservative in terms of providing loans since the financial crisis that occurred around 2009.

2.2 Data cleaning

We now perform some data cleaning. We begin this task by replacing the missing values for both JOB and REASON with “Unknown”, and print out the frequency tables for each of these variables after replacement.

```
levels(dat$JOB) <- c(levels(dat$JOB), "Unknown")
dat$JOB[is.na(dat$JOB)] <- 'Unknown'
levels(dat$REASON) <- c(levels(dat$REASON), 'Unknown')
dat$REASON[is.na(dat$REASON)] <- 'Unknown'
table(dat$JOB)
```

```
##
##      Mgr   Office   Other ProfExe   Sales   Self Unknown
##      767    948    2388    1276    109    193    279
```

```
table(dat$REASON)
```

```
##
## DebtCon HomeImp Unknown
##   3928    1780     252
```

One interesting finding here is that 2,388 individuals had their employment classified as “Other.” When dealing with banks for loans, banks usually direct their attention to an applicant’s employment—and future employment prospects—when deciding whether or not to approve a loan, as they need to know if the applicant will be capable of repaying the loan both now and in the future. Perhaps the real data has more employment classes than the data set we used, but working strictly with the the data that we have on hand, this is an interesting observation.

Next, we perform logarithmic transformations on `LOAN`, `VALUE`, `MORTDUE`, `YOJ`, and `CLAGE`. It should be noted that both `YOJ` and `CLAGE` had some zeros for recordings, thus we used the following to avoid “domain errors,” $\log(x) := \log(x + 1)$ for these variables.

```
dat$LOAN <- log(dat$LOAN)
dat$VALUE <- log(dat$VALUE)
dat$MORTDUE <- log(dat$MORTDUE)
dat$YOJ <- log(dat$YOJ + 1)
dat$CLAGE <- log(dat$CLAGE + 1)
```

Now, we need to impute our missing values, so we have elected to employ `mice()` from the `mice` package, as it is much quicker in terms of time completion relative to the `missForest` package’s imputation.

```
library(mice, quietly=TRUE)
fit.mice <- mice(dat, m=1, maxit=50, method='pmm', seed=500, printFlag=FALSE)
dat.imp <- complete(fit.mice, 1)
```

Finally, we store our known values into a vector—to be used for plotting labels later—and then delete them from the data set, so that we can proceed with the cluster analysis.

```
labs <- dat.imp[, c(1)] # Save for plotting labels later
dat.imp <- dat.imp[, -c(1)] #Remove BAD
```

2.3 Obtaining a distance matrix

We should note that our data set is of mixed types. That is, we have some continuous variables and some dichotomous variables in the same data set. Rather than use `daisy()` with `gower=TRUE`, we will use `model.matrix()` to prepare our data for cluster analysis. This approach has some shortcomings, but we are content with them and their consequences, and thus proceed as follows.

We scale our data set by columns using `scale()` in R. While this may seem reasonable initially, the reader is directed to Hastie, Tibshirani, and Friedman, (2009) for a more in-depth treatment for some problems that may occur with this approach, but we will take their considerations into account and perform this operation none-the-less.

```
library(cluster)
dat.imp <- as.data.frame(dat.imp)
```

```
dat0 <- model.matrix(~.-1, data=dat.imp)
dat0.scaled <- apply(dat0, 2, scale)
```

The methods that we have elected to use for cluster analysis do not require a distance matrix, but only the data, and since the project assignment requires that we compute a distance matrix then we do so here.

```
dat0.dist <- dist(dat0.scaled, method="euclidean")
```

2.4 Clustering algorithms

We have chosen to use K -means and Partitioning Around Medoids (PAM) clustering for this project.¹

2.4.1 K -means clustering

The first method that we use is K -means clustering.

We need to determine the optimal number of clusters to be used, because this “optimal” value must be specified when performing the clustering. While there exists a variety of measures and techniques capable of accomplishing this task, we will only rely on the average silhouette width due to its computational efficiency.

While some readers may wish to rely exclusively on a scree plot and look for an elbow (e.g. the elbow method), this technique is extremely subjective and not much to our liking. Moreover, Everitt and Hothorn (2011), state that the elbow method was initially developed for factor analysis from Cattell (1966), and that caution should be exercised when using this technique for other methods.

Other methods that can be used to determine the optimal number of clusters take much longer to compute than the silhouette width, even when decreasing the number of bootstraps.

```
suppressMessages(library(factoextra))
fviz_nbclust(dat0.scaled, FUNcluster=kmeans, method="silhouette", k.max=20) +
  labs(title= "Optimal number of clusters for K-means")
```

We see that we should select 15 clusters for K -means, from Figure 2. It is worth noting that the width of the silhouette is not too great, and some sources would consider its width as “no substantial structure [having] been found” (Spector, 2007). Nonetheless, we will take this value found to be the optimal number of clusters.

Now, we focus on the arguments for our clustering. James, Witten, Hastie, and Tibshirani, (2013), strongly recommend that `nstart` should be chosen to be between 20 and 50 because

¹One thing that interests me is how statistical methods based on means and medians differ. While PAM is not a strict median—by definition, comparing its performance to a method based on “averages” should prove interesting.




Figure 2: Plot for the average silhouette width used to determine the optimal number of clusters to be used in K-means clustering

“otherwise an undesirable local optimum may be obtained.” So, we choose the minimum of their recommendation. We also set the number of available colors to 15 as this is the maximum number of colors for this clustering method and the one that follows. Moreover, we have tinkered with the parameters for tSNE to achieve a somewhat maximum dispersion as outlined in Wattenberg, Viegas, and Johnson (2016) in Figure 3.

```
library(RColorBrewer)
palette(rainbow(15))
library(Rtsne)
set.seed(5474)
tsne <- Rtsne(dat0.scaled, dims=2, perplexity=50, max_iter=750)
km.cl <- kmeans(dat0.scaled, centers=15, nstart=20)
plot(tsne$Y, t="n", main="tSNE for km")
text(tsne$Y, labels=labs, col=km.cl$cluster)
abline(h=0, v=0, lty=2)
```

Next we print the cluster memberships for our target variable BAD with the following code.

```
table(labs, km.cl$cluster)
```

```
##
## labs   1   2   3   4   5   6   7   8   9  10  11  12  13  14  15
##      0 858 107  71 135 396 738  51 693 268 516 157  95 544  24 118
##      1 144  30  37  57  68 173 144  89  25  73  10   9 215  53  62
```

There are several interesting findings with the output, but these findings will be discussed in the *post hoc* analysis portion of this project.




Figure 3: Plot for K-means clustering with t-SNE.

2.4.2 Partition around medoids (PAM) clustering

The second method that we employ is Partitioning Around Medoids (PAM) clustering. This method has an advantage over K -means as one of the data points must be the center of each cluster; it is more computationally expensive, however.

We need to determine the optimal number of clusters with this method, and for consistency we again examine the silhouette width.

```
fviz_nbclust(dat0.scaled, FUNcluster=pam, method="silhouette", k.max=20) +
labs(title= "Optimal number of clusters for PAM")
```

From Figure 4, we see that we should select 12 clusters, and again the width is fairly low. So, we use `pam()` with `k=12` to find these clusters with the code below, and again plot the memberships with tSNE.

```
plot(tsne$Y, t="n", main="tSNE for pam")
pam.cl <- pam(dat0.scaled, k=12, diss=FALSE)
text(tsne$Y, labels=labs, col=pam.cl$clustering)
abline(h=0, v=0, lty=2)
```

Next we print the cluster memberships for our target variable BAD with the following code.




Figure 4: Plot for the average silhouette width used to determine the optimal number of clusters to be used in PAM clustering.

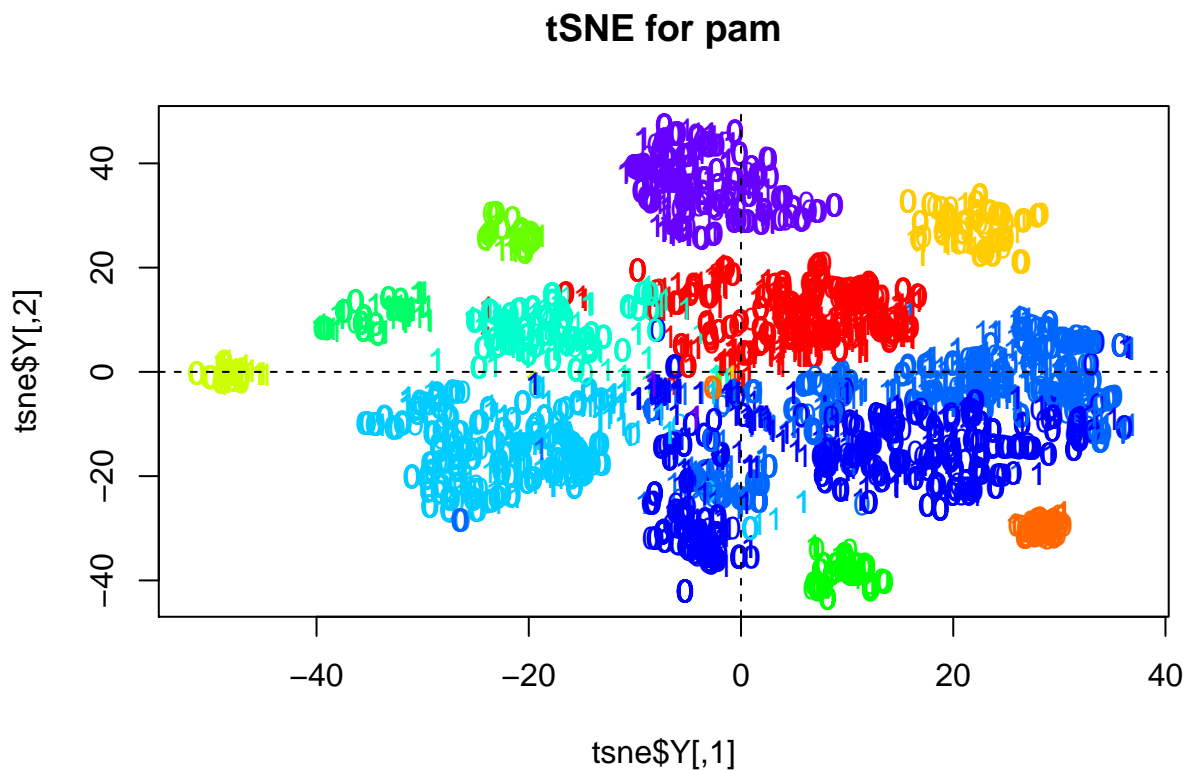


Figure 5: Plot for tSNE with PAM clustering.

```
table(labs, pam.cl$clustering)
```

```
##
## labs    1    2    3    4    5    6    7    8    9   10   11   12
##      0 573   97 268   71 107 159 135 376 724 736 996 529
##      1 249   10  33   38  33  13   58  82 136 233 212  92
```

Looking at the table it is difficult to see if there is a clear difference between the clusters and our target variable BAD. We will look into the cluster memberships for K -means shortly.

2.4.3 Comparing K -means and PAM

Now we compare the clustering results in the following two-way contingency table. We used the Jaccard and Rand similarity indices.

```
library(clusteval)
jac <- cluster_similarity(km.cl$cluster, pam.cl$clustering,
                          similarity="jaccard")
ran <- cluster_similarity(km.cl$cluster, pam.cl$clustering,
                          similarity="rand")
matrix(c("Jaccard", jac, "Rand", ran), byrow=TRUE, ncol=2)
```

```
##      [,1]      [,2]
## [1,] "Jaccard" "0.652755005713718"
## [2,] "Rand"    "0.950016218206965"
```

In the table we see that the two methods are fairly similar. This was somewhat expected, as the methods are similar, but I did not expect the difference of the indices to be as large as it is.

2.5 *Post hoc* analysis

Finally, we perform some *post hoc* analysis using the results from K -means clustering.

We first wish to investigate the relationship of the clusters and the applicant's occupation. The table below shows this relationship.

```
tableJOB <- table(dat$JOB, km.cl$cluster)
tableJOB
```

```
##
##           1    2    3    4    5    6    7    8    9   10   11   12   13   14   15
## Mgr       371  20    0    0   74 110   41    0    0    0    0    0   86   15   50
## Office     0  27    0    0    0    0   34    0 293 589    0    0    0    2    3
## Other     630  67    0    0   14 788   71    0    0    0    0    0 664   38 116
## ProfExe    1  23    0    0  376   13  48 782    0    0    0    0    9   14   10
```

```
## Sales      0  0 108  0  0  0  0  0  0  0  0  0  0  0  1  0
## Self       0  0  0 192  0  0  1  0  0  0  0  0  0  0  0  0
## Unknown    0  0  0  0  0  0  0  0  0  0  0 167 104  0  7  1
```

We see that for most clusters one occupation dominates that cluster. For example, clusters 11 and 12 are made up exclusively with the JOB being unknown, and clusters 3 and 4 only have applicants whose occupation is Sales and Self, respectively.

Next we want to investigate the relationship of the clusters and the applicant's reason for requesting a loan.

```
tableNINQ <- table(dat$REASON, km.cl$cluster)
tableNINQ
```

```
##
##           1    2    3    4    5    6    7    8    9    10   11   12   13
## DebtCon 1002    0   96   73    0  911  158  782    0  589  111    0    0
## HomeImp    0    0   12  114  464    0   37    0  293    0   56    0  759
## Unknown    0  137    0    5    0    0    0    0    0    0    0  104    0
##
##           14   15
## DebtCon   51  155
## HomeImp   20   25
## Unknown    6    0
```

Again, we see that for some clusters the reason for requesting a loan dominates. For example, clusters one, six, eight, and ten only have applicants that are seeking debt consolidation, while clusters 2, 5, 9, and 13 only contain applicants that are seeking funds for home improvement.

Last but not least, we want to investigate the relationship of the clusters and our target variable BAD.

```
tableBAD <- table(dat$BAD, km.cl$cluster)
tableBAD
```

```
##
##      1  2  3  4  5  6  7  8  9  10  11  12  13  14  15
## 0 858 107 71 135 396 738 51 693 268 516 157 95 544 24 118
## 1 144 30 37 57 68 173 144 89 25 73 10 9 215 53 62
```

Here we do not have any clear indications as to whether BAD influenced cluster assignment. Since we do not observe any zeros, we perform a χ^2 test to see if Bad is independent of cluster assignment.

```
chisq.test(tableBAD)
```

```
##
## Pearson's Chi-squared test
##
## data:  tableBAD
```

```
## X-squared = 690.32, df = 14, p-value < 2.2e-16
```

We see that our p -value is essentially zero, and thus the result of this test is significant for any reasonable level of significance α . So, we conclude that BAD is not independent of cluster membership.

3 Discussion

Due to the large number of clusters in both K -means and PAM, the coloring for the tSNE plots is not ideal. I have experimented with different color schemes, but ultimately settled for `rainbow(15)`. I say that it is not ideal because the plots are not colorblind friendly, and it can be difficult to separate several shades of blue and green, for those not color blind.

I also need to make a statement regarding academic honesty. I would not have been able to complete the proof for this project without a hint from Zheng Zhang. I also helped Simon (Xiang) Teng with some of his program code. He wrote the code himself, but I merely provided some tips. I am the Statistical Programming course's Graduate Teaching Assistant—something that I absolutely love, but wish more students would visit me during office my hours—thus I would like to see him succeed in his R programming endeavours. Fortunately, they are beginning Python (my native programming language), so this issue should not present itself for the remainder of this semester and future projects.

4 References

- [0] Cattell, R. B. (1966), "The scree test for the number of factors," *Multivariate Behavioural Research*, 1, 245–276.
- [1] Everitt, B., Hothorn, T. (2011), *An Introduction to Applied Multivariate Analysis with R*, New York: Springer.
- [2] Hastie, T., Tibshirani, R., and Freidman, J. (2009), *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, New York: Springer
- [3] James, G., Witten, D., Hastie, T., and Tibshirani, R. (2013), *An Introduction to Statistical Learning with Applications in R*, New York: Springer.
- [4] Spector, P. (2007), "Performing and interpreting cluster analysis," Retrieved from: <https://www.stat.berkeley.edu/~spector/s133/Clus.html> .
- [5] Wattenberg, M., Viégas, F., Johnson, I. (2016), "How to use t-SNE effectively", *Distill*, doi:10.23915/distill.00002.