

Project Seven for Data Mining

Raymond Anthony Ford
raford2@miners.utep.edu

Due: 14 May 2018

Contents

1	Data preparation	1
2	Supervised learning	2
2.1	Linear discriminant analysis (LDA)	2
2.2	Logistic regression with backward elimination	2
2.3	Random Forests	2
2.4	Model comparison	3
3	Additional features from Random Forests	3
A	Appendix	8

1 Data preparation

We begin this project by bringing the data into R. The data comes from the `kernlab` package, and we seek to use several classification methods to determine if an email is a regular email or a spam email.

```
library(kernlab)
data(spam)
dat <- spam
```

The data set contains 4601 emails in total and of which 1813 were spam emails, so approximately 40% of the emails in the data set are spam. Looking at the data we see that the type of each variable, except for our target variable, is continuous. Our target variable `type` is a character variable that is recorded as either spam or nonspam. Moreover, we note that there are no missing values in our dataset.¹

We next split the data into a training and test set with the following code.

```
set.seed(5474)
training.data.index <- sample(1:nrow(dat), 0.667*nrow(dat))
```

¹See the appendix for the ‘R’ code and output used for this portion.

```
training <- dat[training.data.index, ]  
test <- dat[-training.data.index, ]
```

2 Supervised learning

We next fit models using three different statistical learning methods to classify an email as either spam or nonspam; these methods are listed in each of the subsections below.

2.1 Linear discriminant analysis (LDA)

Our first method is linear discriminate analysis (LDA). The code below is used to fit said model.

```
library(MASS)  
spam.lda <- lda(type ~ ., data=training, CV=FALSE)  
spam.lda.pred <- as.vector(predict(spam.lda, test)$x)
```

2.2 Logistic regression with backward elimination

Our second method is a logistic regression model selected via backward elimination based on BIC. The code below is used to fit this model.

```
fit.full <- suppressWarnings(glm(type~., data=training, family = binomial))  
spam.back <- suppressWarnings(stepAIC(fit.full, direction = "backward",  
                                     k=log(nrow(dat)), trace=FALSE))  
spam.back.pred <- predict(spam.back, newdata=test, type="response")
```

2.3 Random Forests

Our final classification model is a random forest. We have elected to use the default values as arguments/parameters when fitting this model. The code for this is below.

```
suppressMessages(library(randomForest, quietly=TRUE))  
set.seed(5474)  
spam.rf <- randomForest(type~., data=training)  
spam.rf.pred <- as.data.frame(predict(spam.rf, test, type="prob"))$spam
```




Figure 1: ROC(AUC) curves for each of the classification methods employed: linear discriminant analysis (LDA), logistic regression, and random forests, respectively.

2.4 Model comparison

Finally, we evaluate the performance of these three models by comparing their area under the receiver operating characteristic curve (AUC) values.

```
suppressMessages(suppressWarnings(library(pROC, quietly=TRUE)))
par(mfrow=c(1, 3), mar=rep(4,4), pty="s")
plot.roc(test$type, spam.lda.pred, main="(a) LDA", percent=TRUE,
  print.auc=TRUE, print.auc.cex=1.0, col="blue")
plot.roc(test$type, spam.back.pred, main="(b) Logistic Regression", percent=TRUE,
  print.auc=TRUE, print.auc.cex=1.0, col="blue")
plot.roc(test$type, spam.rf.pred, main="(c) Random Forests", percent=TRUE,
  print.auc=TRUE, print.auc.cex=1.0, col="blue")
```

From Figure 1, we see that the random forests method out performs the other methods with an AUC of 98.7%.

3 Additional features from Random Forests

Finally, we train a random forest model with 2000 trees using the entire data set.

```
set.seed(5474)
spam.rf.full <- randomForest(type~., data=dat, importance=TRUE, proximity=TRUE,
  ntree=2000)
varImpPlot(spam.rf.full, main="Variable importance")
```

From the variable importance plots in Figure 2, we see that `charExclamation` and `charDollar` (the symbols “!” and “\$”, respectively) are the top two most important variables according to the mean decrease in Gini index. We next plot the partial dependence plots for these two variables.

Variable importance

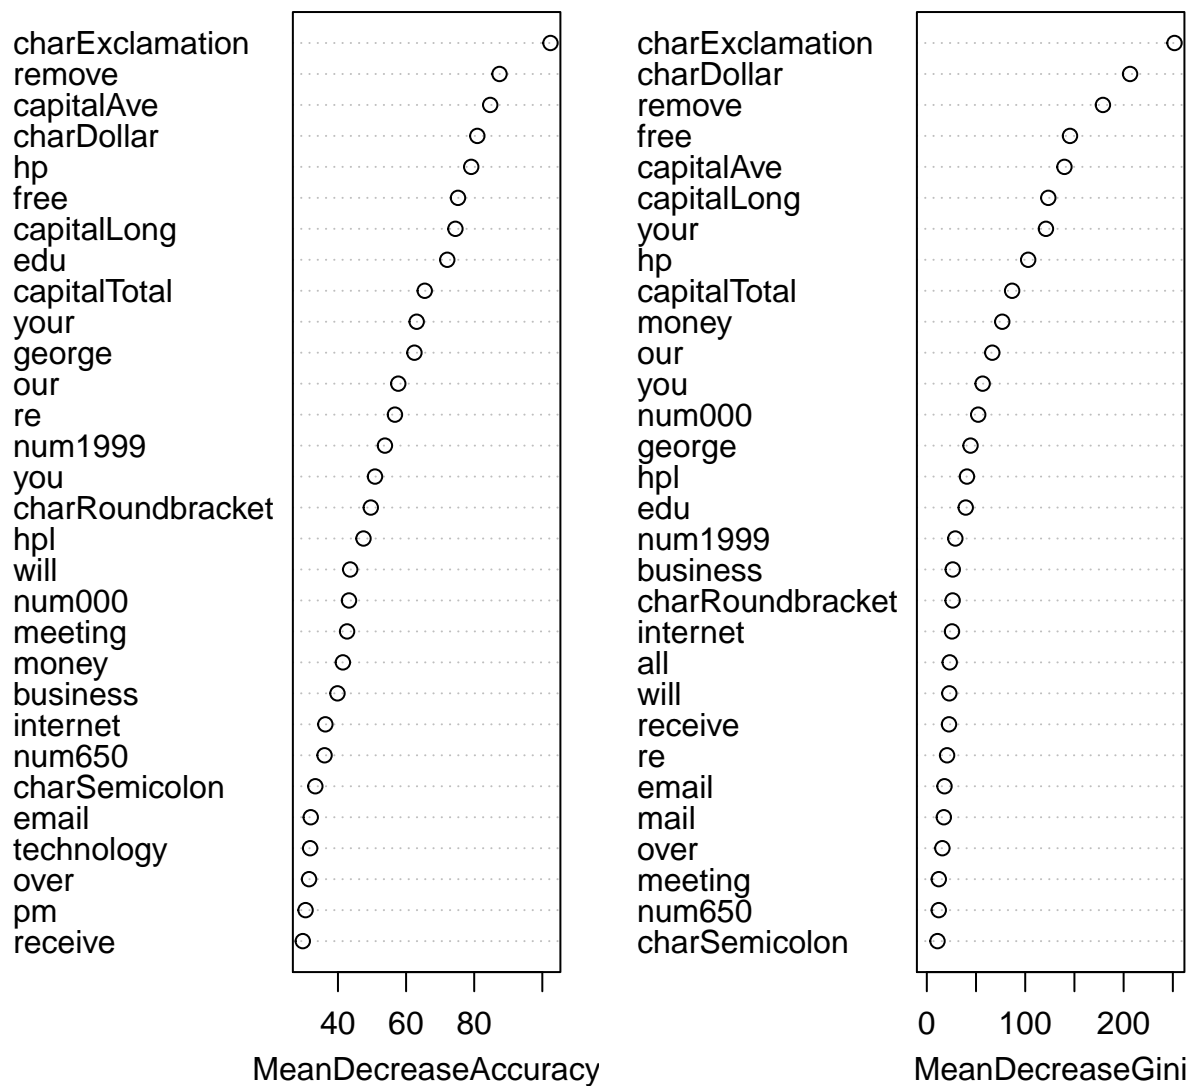


Figure 2: Variable importance plot for the random forest fit with the full data set.

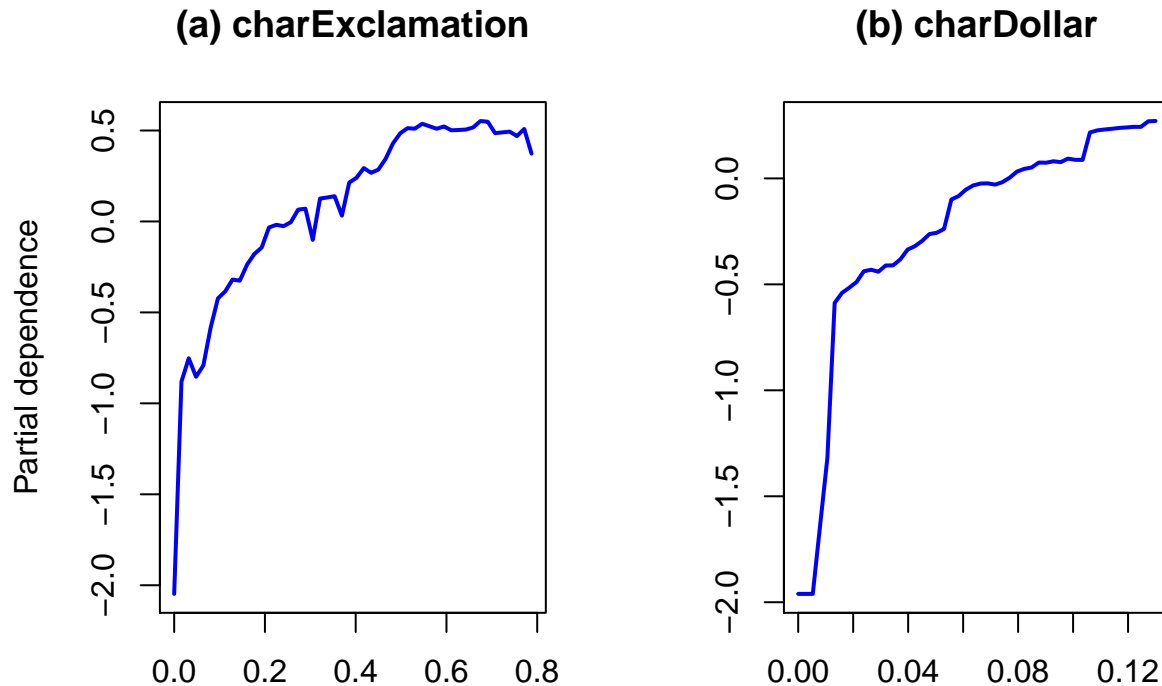


Figure 3: Partial dependence plots for top two important variables.

```
suppressMessages(library(interpretR, quietly=TRUE))
par(mfrow=c(1, 2))
parDepPlot(x.name="charExclamation", spam.rf.full, data=dat,
           main="(a) charExclamation", col="blue", lwd=2, xlab="",
           ylab="Partial dependence")
parDepPlot(x.name="charDollar", spam.rf.full, data=dat,
           main="(b) charDollar", col="blue", lwd=2, xlab="", ylab="")

library(beepR);beep(8)
```

In the partial dependence plots in Figure 3, we see that as the number of exclamation marks (a) and dollar signs (b) increase,² the more likely an email will be classified as a spam email.

Lastly, we transform the proximity matrix into a dissimilarity matrix and plot the first two principal components.

```
DIST <- 1 - spam.rf.full$proximity
labs <- dat[,58]
set.seed(5474)
pca.res <- prcomp(dat[, -58], retx=TRUE)
plot(pca.res$x[,1:2], pch="", main="PC1 and PC2 for email spam data set")
text(pca.res$x[,1:2], col=c("blue", "orange")[labs], lab=labs)
abline(h=0, v=0, lty=2)
```

²Outliers are removed in these plots.




Figure 4: Plot of the spam data set using PCA.

In the above plot, we see that there appears to be a clear clustering of spam and nonspam emails. Moreover, there are many outlying values for the PCs for spam emails, but when these outlying PCs are removed the pattern remains the same.³

³I have omitted the plot with the outlying PCs removed, but if needed I can email it to you.

A Appendix

Words

```
n <- nrow(dat)
out <- NULL
for (k in 1:ncol(dat)) {
  vname <- colnames(dat)[k]
  x <- as.vector(dat[,k])
  n1 <- sum(is.na(x), na.rm=TRUE)
  n2 <- sum(x=="NA", na.rm=TRUE)
  n3 <- sum(x=="'", na.rm=TRUE)
  n4 <- sum(x=="?", na.rm=TRUE)
  n5 <- sum(x=="*", na.rm=TRUE)
  n6 <- sum(x==".", na.rm=TRUE)
  nmiss <- n1 + n2 + n3 + n4 + n5 + n6
  ncomplete <- n - nmiss
  var.type <- typeof(x)
  if (var.type == "integer") {
    if (length(unique(x)) == 2) {
      out <- rbind(out, c(col.number=k, vname=vname, mode="binary",
                          n.levels=length(unique(x)), ncomplete=ncomplete,
                          miss.prop=round(nmiss/n, digits=4)))
    } else {
      out <- rbind(out, c(col.number=k, vname=vname, mode=typeof(x),
                          n.levels=length(unique(x)), ncomplete=ncomplete,
                          miss.prop=round(nmiss/n, digits=4)))
    }
  } else {
    out <- rbind(out, c(col.number=k, vname=vname, mode=typeof(x),
                        n.levels=length(unique(x)), ncomplete=ncomplete,
                        miss.prop=round(nmiss/n, digits=4)))
  }
}
out <- as.data.frame(out)
row.names(out) <- NULL
out
```

##	col.number	vname	mode	n.levels	ncomplete	miss.prop
## 1	1	make	double	142	4601	0
## 2	2	address	double	171	4601	0
## 3	3	all	double	214	4601	0
## 4	4	num3d	double	43	4601	0
## 5	5	our	double	255	4601	0
## 6	6	over	double	141	4601	0

## 7	7	remove	double	173	4601	0
## 8	8	internet	double	170	4601	0
## 9	9	order	double	144	4601	0
## 10	10	mail	double	245	4601	0
## 11	11	receive	double	113	4601	0
## 12	12	will	double	316	4601	0
## 13	13	people	double	158	4601	0
## 14	14	report	double	133	4601	0
## 15	15	addresses	double	118	4601	0
## 16	16	free	double	253	4601	0
## 17	17	business	double	197	4601	0
## 18	18	email	double	229	4601	0
## 19	19	you	double	575	4601	0
## 20	20	credit	double	148	4601	0
## 21	21	your	double	401	4601	0
## 22	22	font	double	99	4601	0
## 23	23	num000	double	164	4601	0
## 24	24	money	double	143	4601	0
## 25	25	hp	double	395	4601	0
## 26	26	hpl	double	281	4601	0
## 27	27	george	double	240	4601	0
## 28	28	num650	double	200	4601	0
## 29	29	lab	double	156	4601	0
## 30	30	labs	double	179	4601	0
## 31	31	telnet	double	128	4601	0
## 32	32	num857	double	106	4601	0
## 33	33	data	double	184	4601	0
## 34	34	num415	double	110	4601	0
## 35	35	num85	double	177	4601	0
## 36	36	technology	double	159	4601	0
## 37	37	num1999	double	188	4601	0
## 38	38	parts	double	53	4601	0
## 39	39	pm	double	163	4601	0
## 40	40	direct	double	125	4601	0
## 41	41	cs	double	108	4601	0
## 42	42	meeting	double	186	4601	0
## 43	43	original	double	136	4601	0
## 44	44	project	double	160	4601	0
## 45	45	re	double	230	4601	0
## 46	46	edu	double	227	4601	0
## 47	47	table	double	38	4601	0
## 48	48	conference	double	106	4601	0
## 49	49	charSemicolon	double	313	4601	0
## 50	50	charRoundbracket	double	641	4601	0
## 51	51	charSquarebracket	double	225	4601	0

## 52	52	charExclamation	double	964	4601	0
## 53	53	charDollar	double	504	4601	0
## 54	54	charHash	double	316	4601	0
## 55	55	capitalAve	double	2161	4601	0
## 56	56	capitalLong	double	271	4601	0
## 57	57	capitalTotal	double	919	4601	0
## 58	58	type character		2	4601	0